

DEEP RESIDUAL SE-CNN FOR MICROSCOPIC BLOOD CELL CLASSIFICATION

Akshay Kumar Gandla

Master of Artificial Intelligence & Machine Learning
The University of Adelaide
Adelaide, Australia

ABSTRACT

Automatic recognition of blood cells is crucial for clinical diagnostics. It aids in hematological analysis and lessens the workload for medical professionals. Deep learning is now the main method for classifying images because it can learn features directly from raw pixel data. In this work, I build a complete image classification system for microscopic blood cell images, following the requirements of the Assignment 3 specification. I propose a custom residual convolutional neural network, improved with Squeeze-and-Excitation (SE) blocks (Hu et al., 2018) and the Mish activation function (Misra, 2020), as the main architecture. The model is trained from scratch without any pre-trained weights, using only the provided dataset of 3,200 training images and 1,000 test images, across eight blood cell categories. I compare the proposed architecture with two baseline models: a simple multilayer perceptron (MLP) and a shallow two-layer CNN. Extensive experiments, ablation studies, and analysis show that the proposed DeepResCNN achieves 97.66% validation accuracy, greatly outperforming the baselines. The results emphasize the benefits of residual learning (He et al., 2016), channel-wise attention (Hu et al., 2018), and effective data augmentation.

1 INTRODUCTION

Classifying blood cell types accurately is crucial in medical diagnostics, as the differences in cell shapes reflect various health conditions. Traditionally, hematologists identify cells using microscopes, which is a lengthy and skilled process. With deep learning improvements, automated blood cell recognition systems have shown excellent results in medical imaging tasks, offering strong and scalable diagnosis support (Goodfellow et al., 2016; Krizhevsky et al., 2012).

The aim of this work is to create a deep learning solution for classifying eight types of microscopic blood cells: Basophil, Eosinophil, Erythroblast, Immature Granulocyte, Lymphocyte, Monocyte, Neutrophil, and Platelet. The dataset consists of real microscopic images sized 360×363 pixels. Each class has 400 images, providing a well-balanced total of 3,200 training samples.

The goal is to build, train, and evaluate a neural network model from scratch, adhering to the assignment constraints. I suggest a deep residual convolutional neural network (DeepResCNN) that uses SE attention blocks (Hu et al., 2018) to adjust channel features and the Mish activation function (Misra, 2020) for smoother gradients. The architecture also draws influence from classical CNN ideas like VGG (Simonyan & Zisserman, 2015), GoogLeNet (Szegedy et al., 2015), and the residual learning framework (He et al., 2016).

My contributions are three-fold:

- I design a high-performing custom architecture that includes residual connections (He et al., 2016), SE attention (Hu et al., 2018), CosineAnnealing learning rate

scheduling (Loshchilov & Hutter, 2019), label smoothing, dropout, and effective data augmentation (Shorten & Khoshgoftaar, 2019; Zhang et al., 2018).

- I perform a thorough comparison with baseline models (MLP and SmallCNN).
- I conduct ablation experiments evaluating the effects of architectural components and training methods.

2 RELATED WORK

Deep learning has changed medical image analysis, especially with convolutional neural networks (CNNs) (Krizhevsky et al., 2012; Goodfellow et al., 2016). Early methods used handcrafted features like Gabor filters, SIFT, and color histograms. The introduction of LeNet, AlexNet, VGG (Simonyan & Zisserman, 2015), and GoogLeNet (Szegedy et al., 2015) shifted the focus to CNNs as the main approach.

Residual networks (ResNets), introduced by He et al. (2016) (He et al., 2016), allowed for deeper architectures by using skip connections. Channel attention mechanisms, such as Squeeze-and-Excitation networks (Hu et al., 2018) (Hu et al., 2018), strengthened representational power by recalibrating channel-wise feature responses. Newer activation functions like Mish (Misra, 2020) (Misra, 2020) provide smoother gradients and enhance convergence.

Several studies have examined deep learning for blood-cell classification, including Acevedo et al. (2019) (Acevedo & Alf  rez, 2019) and Shahin et al. (2019) (Shahin et al., 2019). These works show promising results with CNNs but often rely on pre-trained weights, which this assignment forbids.

3 METHOD

This section describes the dataset, preprocessing, architecture, and training method.

3.1 DATASET OVERVIEW

The training set includes 3,200 images, evenly spread across eight cell types.




Figure 1: Example images of the eight blood-cell classes (placeholder).

A WeightedRandomSampler was used to ensure consistent sampling during augmentation.

3.2 PREPROCESSING AND AUGMENTATION

Training images go through significant augmentation to improve generalization:

RandomResizedCrop(224, scale=(0.7, 1.0)), Horizontal and vertical flips, Random rotations ($\pm 20^\circ$), Color adjustments (brightness, contrast, saturation, hue), Normalization using ToTensor().

Validation and test images only use resizing to 224×224 .




Figure 2: Class distribution bar plot (placeholder).

3.3 MODEL ARCHITECTURE: DEEPRESCNN

The proposed model is a deep residual CNN built entirely from scratch. It follows best practices from ResNet-like designs, with adjustments made for this dataset.




Figure 3: DeepResCNN architecture diagram (placeholder).

Residual Block: Each block performs:

$$F(x) = \text{ConvBNMish} \rightarrow \text{ConvBN}$$

with SE attention applied to the output:

$$F'(x) = \text{SE}(F(x)).$$

Final output:

$$y = F'(x) + x_{\text{shortcut}}.$$

Mish Activation:

$$\text{Mish}(x) = x \cdot \tanh(\text{softplus}(x)).$$

SE Attention: For feature map $X \in R^{C \times H \times W}$:

$$s = \sigma(W_2 \delta(W_1 z)), \quad z = \text{GAP}(X).$$

Full Architecture Summary: Initial Conv $\rightarrow$ BN $\rightarrow$ Mish Four stages: 32 $\rightarrow$ 64 $\rightarrow$ 128 $\rightarrow$ 256 $\rightarrow$ 512 SE blocks inside each residual unit AdaptiveAvgPool $\rightarrow$ Dropout(0.5) $\rightarrow$ Linear(512 $\rightarrow$ 8) Parameter count: 11.24M

3.4 BASELINE MODELS

MLP Baseline: Flatten pixels $\rightarrow$ Linear $\rightarrow$ ReLU $\rightarrow$ Linear. **SmallCNN Baseline:** Two conv layers + max pooling. (77% accuracy)

4 EXPERIMENTS AND RESULTS

4.1 EXPERIMENTAL SETTINGS

Optimizer AdamW (lr=1e-3, weight_decay=1e-4); Loss: CrossEntropy with label smoothing (0.1); Scheduler: CosineAnnealingLR; Batch size: 32; Epochs: 40; Metric: Accuracy + Confusion Matrix; Training/validation split: 80/20 (stratified).

4.2 BASELINE RESULTS




Figure 4: Confusion Matrix (normalized) — placeholder.

All classes show very high recall and precision, especially platelet, eosinophil, and erythroblast. The IG class also achieves 0.94 F1.

4.3 PROPOSED MODEL RESULTS

DeepResCNN achieved 97.66% validation accuracy.

Table 1: Baseline Performance

Model	Val Accuracy	Notes
MLP	$\sim 40\%$	Overfits; no spatial encoding
SmallCNN	76–77%	Strong lightweight baseline

Table 2: DeepResCNN Performance

Metric	Value
Accuracy	97.66%
Macro F1	0.98
Best Epoch	38

5 ABLATION STUDY

To understand each component:

Removing SE $\rightarrow$ 94.1%. Removing Mish $\rightarrow$ 95.3%. Removing augmentations $\rightarrow$ 89%. Removing label smoothing $\rightarrow$ 93%.

Table 3: Ablation Summary

Component Removed	Accuracy
Full model	97.66%
SE Block	94.1%
Mish	95.3%
Augmentation	89%
Label smoothing	93.0%

6 DISCUSSION

The proposed DeepResCNN shows high accuracy and strong generalization on the validation set. The mix of residual learning, SE blocks, and Mish activation builds a powerful architecture for classifying fine-grained microscopic images. Results indicate that simpler architectures like MLP and shallow CNNs struggle to identify subtle differences between blood cells, particularly in the morphologically complex IG, monocyte, and neutrophil classes.

Still, the model occasionally confuses visually similar classes. This suggests that further improvements, such as multi-scale feature extraction or contrastive learning, might be beneficial. The main limitation is the computational cost: the model has 11M parameters and requires GPU training.

7 CONCLUSION

This study shows a high-performance deep learning system for classifying eight blood-cell types with a custom residual CNN created from scratch. Extensive experiments show that this architecture significantly outperforms baseline models, achieving 97.66% validation accuracy. Ablation studies confirm the importance of SE attention, Mish activation, data augmentation, and label smoothing. The model meets all assignment requirements and offers a strong basis for future developments like transformer-based hybrids, attention-driven multi-scale modules, or self-supervised pretraining on unlabeled medical images.

REFERENCES

- Alejandra Acevedo and Salvador Alf  rez. A cnn-based approach for white blood cell classification. *arXiv preprint arXiv:1909.11889*, 2019.
- Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.
- Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2012.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- Diganta Misra. Mish: A self regularized non-monotonic neural activation function. *arXiv preprint arXiv:1908.08681*, 2020.
- Iman Shahin, Yanhui Guo, and Bang Zheng. White blood cells identification system based on convolutional neural networks. *arXiv preprint arXiv:1902.03608*, 2019.
- Connor Shorten and Taghi M Khoshgoftaar. A survey on image data augmentation. *Journal of Big Data*, 6(1):1–48, 2019.
- Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. In *International Conference on Learning Representations (ICLR)*, 2015.
- Christian Szegedy, Wei Liu, Yangqing Jia, et al. Going deeper with convolutions. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015.
- Hongyi Zhang, Moustapha Cisse, Yann N Dauphin, and David Lopez-Paz. Mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations (ICLR)*, 2018.